

# Data102 Final Project Elections

December 12, 2021

## 1 Group Members

1. Sarah Shikanov, sarahshikanov@berkeley.edu
2. James Marquez, jamesamarquez@berkeley.edu
3. Michael Chien, mchien512@berkeley.edu
4. Rachel Fei, wfei@berkeley.edu

```
[1]: import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
import arviz as az
import geopandas as gp
```

### 1.1 Load in the Data

```
[2]: dem_df = pd.read_csv('dem_candidates.csv', encoding='latin1')
rep_df = pd.read_csv('rep_candidates.csv', encoding='latin-1')
```

```
[3]: poll_data = pd.read_csv('presidential_polls.csv')
```

1.1.1 Dataset was taken from <https://www.kaggle.com/fivethirtyeight/2016-election-polls>

## 2 1.0 Data Overview

The democrat and republican data were generated by sampling. This was taken from the class dataset. At quick glance from the our pre processing step below we can see that in our first entry we have a governor running from office for Alabama. At quick glance we can see that the Partisan Lean is heavily negative which means its heavily red. This is consistent with what we would expect from before we began doing anything. We are not worried much about bias as this data is simply a representation of endorsements and the primary. Important features that would be useful if there was a way to get the total contribution with the candidate. We tried merging the dataframes with another dataset but it was very non trivial and thus decided to work with what we had. This dataset is public so all participants are aware of this dataset. As to the granularity each row represents a candidate running for office with endorsements and primary election data.

The 2016 Polling dataset was also generated by sampling. This was a dataset from Kaggle and we

decided to use this because the election dataset given only covered endorsements and there weren't very many cool problems we could think of analyzing. Again by looking at the data we can get a quick glance at the polling advantages for each candidate and they seem to be consistent. The participants are aware of the collection and use of the data as it's published by pollsters. We can see Unlike the class dataset the polling dataset is worrisome with selection bias due to the unreliability of some pollsters but there is a column for grade which should help us identify good pollsters. The columns are sufficient in this dataset to perform what we want.

## 2.1 1.1 Preprocessing.

```
[62]: dem_df['Party'] = "Democrat"
      rep_df['Party'] = 'Republican'
```

```
[63]: df = pd.concat([dem_df, rep_df])
      df
```

```
[63]:
```

	Candidate	State	District	Office	Type	Race	Type	\
0	Anthony White	(Alabama)	AL	Governor of Alabama	Governor	Regular		
1	Christopher Countryman	AL	Governor of Alabama	Governor	Regular			
2	Doug "New Blue" Smith	AL	Governor of Alabama	Governor	Regular			
3	James C. Fields	AL	Governor of Alabama	Governor	Regular			
4	Sue Bell Cobb	AL	Governor of Alabama	Governor	Regular			
..	...	...	...	...	...	...		
769	Bill Dahlin	WY	Governor of Wyoming	Governor	Regular			
770	Harriet Hageman	WY	Governor of Wyoming	Governor	Regular			
771	Sam Galeotos	WY	Governor of Wyoming	Governor	Regular			
772	Foster Friess	WY	Governor of Wyoming	Governor	Regular			
773	Taylor Haynes	WY	Governor of Wyoming	Governor	Regular			

	Race	Primary Election Date	Primary Status	Primary Runoff Status	\
0		6/5/18	Lost	None	
1		6/5/18	Lost	None	
2		6/5/18	Lost	None	
3		6/5/18	Lost	None	
4		6/5/18	Lost	None	
..		...	...	...	
769		8/21/18	Lost	None	
770		8/21/18	Lost	None	
771		8/21/18	Lost	None	
772		8/21/18	Lost	None	
773		8/21/18	Lost	None	

	General Status	Partisan Lean	...	Great America Endorsed?	NRA Endorsed?	\
0	None	-28.879999	...	NaN	NaN	
1	None	-28.879999	...	NaN	NaN	
2	None	-28.879999	...	NaN	NaN	
3	None	-28.879999	...	NaN	NaN	

4	None	-28.879999	...		NaN	NaN
..	...	...	...	...	...	...
769	None	NaN	...		NaN	NaN
770	None	NaN	...		NaN	NaN
771	None	NaN	...		NaN	NaN
772	None	NaN	...		NaN	NaN
773	None	NaN	...		NaN	NaN

	Right to Life Endorsed?	Susan B. Anthony Endorsed?	\
0	NaN	NaN	
1	NaN	NaN	
2	NaN	NaN	
3	NaN	NaN	
4	NaN	NaN	
..	...	...	
769	No	NaN	
770	Yes	NaN	
771	Yes	NaN	
772	Yes	NaN	
773	Yes	NaN	

	Club for Growth Endorsed?	Koch Support?	House Freedom Support?	\
0	NaN	NaN	NaN	
1	NaN	NaN	NaN	
2	NaN	NaN	NaN	
3	NaN	NaN	NaN	
4	NaN	NaN	NaN	
..	...	...	...	
769	NaN	NaN	NaN	
770	NaN	NaN	NaN	
771	NaN	NaN	NaN	
772	NaN	NaN	NaN	
773	NaN	NaN	NaN	

	Tea Party Endorsed?	Main Street Endorsed?	Chamber Endorsed?
0	NaN	NaN	NaN
1	NaN	NaN	NaN
2	NaN	NaN	NaN
3	NaN	NaN	NaN
4	NaN	NaN	NaN
..	...	...	...
769	NaN	NaN	NaN
770	NaN	NaN	NaN
771	NaN	NaN	NaN
772	NaN	NaN	NaN
773	NaN	NaN	NaN

[1563 rows x 46 columns]

```
[6]: # Democrats have weird NaN values for candidates who advanced as well as
      ↪ candidates who lost so we drop these values
      # there are 22 of them.
      # Republicans have 0 so we don't have to do anything.
      bad_indices = list(dem_df[(dem_df['Won Primary'].isna())].index)
      dem_df = dem_df.drop(bad_indices);
```

```
[7]: # Treat a Nan Value as 0 and then encode no and yes as 0 and 1s.
      df.fillna(0, inplace=True)
      df.replace('No', 0, inplace=True)
      df.replace('Yes', 1, inplace=True)
```

```
[8]: # Get a list of all democrat and republican endorsements
      dem_endorsements = list(df.columns[13:32])
      rep_endorsements = list(df.columns[33:46])
      endorsements = dem_endorsements + rep_endorsements
```

## 3 2.0 EDA

### 3.1 2.1 EDA on 2018 Primary Election Dataset

Lets begin with some EDA on our dataest with republicans and democrats from five thirty eight.

```
[9]: dem_endorsements
```

```
[9]: ['Veteran?',
      'LGBTQ?',
      'Elected Official?',
      'Self-Funder?',
      'STEM?',
      'Obama Alum?',
      'Party Support?',
      'Emily Endorsed?',
      'Guns Sense Candidate?',
      'Biden Endorsed?',
      'Warren Endorsed? ',
      'Sanders Endorsed?',
      'Our Revolution Endorsed?',
      'Justice Dems Endorsed?',
      'PCCC Endorsed?',
      'Indivisible Endorsed?',
      'WFP Endorsed?',
      'VoteVets Endorsed?',
      'No Labels Support?']
```

We have a prior belief that some endorsements matter more than others. For example we take that Party Support should have a big impact on the number of votes one receives. In addition we think that big names such as Sanders and Biden.

```
[10]: dem_df.fillna(0, inplace=True)
dem_df.replace('No', 0, inplace=True)
dem_df.replace('Yes', 1, inplace=True)
```

```
[11]: dem_endorsements_df = dem_df[['Candidate'] + dem_endorsements + ['Won Primary']]
dem_endorsements_df
```

```
[11]:
```

	Candidate	Veteran?	LGBTQ?	Elected Official?	\
0	Anthony White (Alabama)	1	0	0	
1	Christopher Countryman	0	1	0	
2	Doug "New Blue" Smith	1	0	0	
3	James C. Fields	1	0	1	
4	Sue Bell Cobb	0	0	1	
..	...	...	...	...	
806	Talley Sergeant	0	0	0	
807	Janice Hagerman	0	0	0	
808	Paul Davis	0	0	0	
809	Richard Ojeda	1	0	1	
810	Shirley Love	0	0	0	

	Self-Funder?	STEM?	Obama Alum?	Party Support?	Emily Endorsed?	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	
..	...	...	...	...	...	
806	0	0	1	0	0	
807	0	0	0	0	0	
808	0	0	0	0	0	
809	0	0	0	0	0	
810	0	0	0	0	0	

	Guns Sense Candidate?	...	Warren Endorsed?	Sanders Endorsed?	\
0	0	...	0	0	
1	0	...	0	0	
2	0	...	0	0	
3	0	...	0	0	
4	0	...	0	0	
..	...	...	...	...	
806	0	...	0	0	
807	0	...	0	0	
808	0	...	0	0	

809	0	...	0	0
810	0	...	0	0

	Our Revolution Endorsed?	Justice Dems Endorsed?	PCCC Endorsed?	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
..	...	...	...	
806	0	0	0	
807	0	0	0	
808	0	0	0	
809	0	0	0	
810	0	0	0	

	Indivisible Endorsed?	WFP Endorsed?	VoteVets Endorsed?	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
..	...	...	...	
806	0	0	0	
807	0	0	0	
808	0	0	0	
809	0	1	1	
810	0	0	0	

	No Labels Support?	Won Primary
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
..	...	...
806	0	1
807	0	0
808	0	0
809	0	1
810	0	0

[789 rows x 21 columns]

```
[12]: dem_copy = dem_df.copy()
dem_copy = dem_df[['Party Support?', 'Biden Endorsed?', 'Sanders Endorsed?'] +
↳ ['Primary %']]
```

```
dem_copy
```

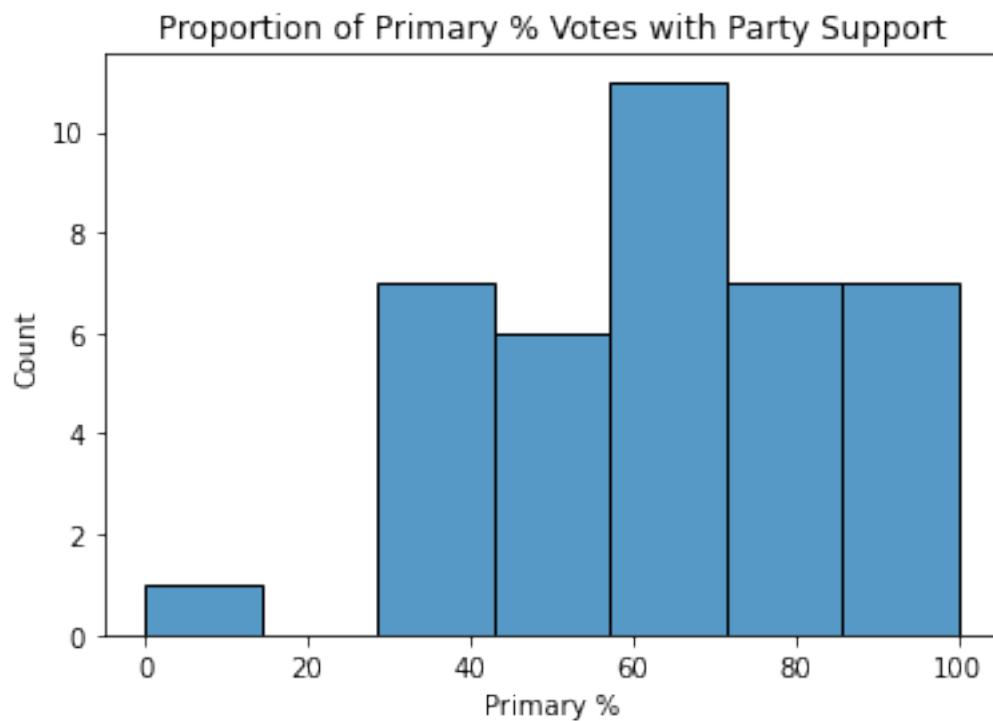
```
[12]:
```

	Party Support?	Biden Endorsed?	Sanders Endorsed?	Primary %
0	0	0	0	3.420000
1	0	0	0	1.740000
2	0	0	0	3.270000
3	0	0	0	8.000000
4	0	0	0	28.980000
..	...	...	...	...
806	0	0	0	62.570000
807	0	0	0	7.240000
808	0	0	0	15.960000
809	0	0	0	52.160000
810	0	0	0	24.639999

```
[789 rows x 4 columns]
```

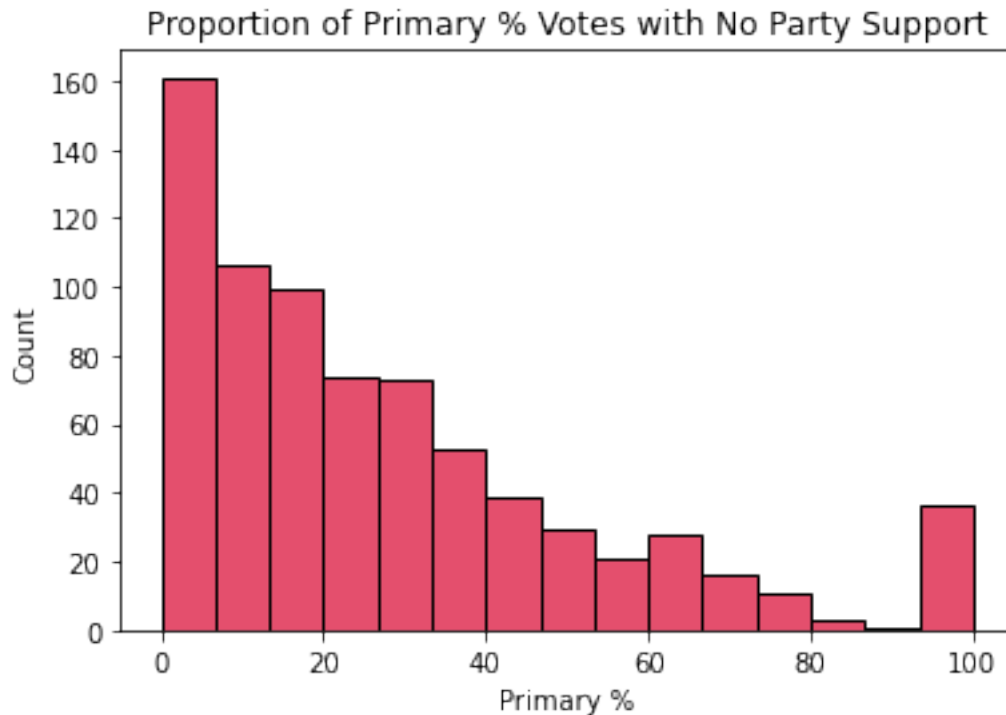
```
[13]: party_support = dem_copy[dem_copy['Party Support?'] == 1]
sns.histplot(party_support['Primary %'])
plt.title('Proportion of Primary % Votes with Party Support')
party_support.shape
```

```
[13]: (39, 4)
```



```
[14]: no_party_support = dem_copy[dem_copy['Party Support?'] == 0]
sns.histplot(no_party_support['Primary %'], color='crimson')
plt.title('Proportion of Primary % Votes with No Party Support')
no_party_support.shape
```

```
[14]: (750, 4)
```



From here we can see that party support is quite important. The histogram with no party support is right skewed with many of the votes not even above 50%. Out of 750 candidates there is a few percentage that is above 50% compared to the first histogram. We should keep this in mind when building our model.

Let's see if we can get a good view of partisan lean based on the state.

```
[26]: df_copy = df.copy()
df_copy = df[['Partisan Lean', 'State']]
df_copy = df_copy.groupby(by=['State']).mean()
df_copy.reset_index().head(5)
```

```
[26]: State  Partisan Lean
0    AK      0.000000
1    AL    -36.023125
2    AR    -23.818889
3    AZ     -1.637333
```



```
4    CA    -0.819133
```

```
[27]: map_us.head(5)
```

```
[27]: STATEFP  STATENS  AFFGEOID GEOID STUSPS  NAME LSAD \
0      06  01779778  0400000US06    06    CA      California  00
1      11  01702382  0400000US11    11    DC  District of Columbia  00
2      12  00294478  0400000US12    12    FL      Florida  00
3      13  01705317  0400000US13    13    GA      Georgia  00
4      16  01779783  0400000US16    16    ID      Idaho  00
```

```
      ALAND  AWATER  region \
0  403483823181  20483271881    West
1    158350578    18633500  Northeast
2  138903200855  31407883551  Southeast
3  148963503399  4947080103  Southeast
4  214045425549  2397728105    West
```

```
      geometry
0  MULTIPOLYGON Z (((-118.59397 33.46720 0.00000,...
1  POLYGON Z ((-77.11976 38.93434 0.00000, -77.04...
2  MULTIPOLYGON Z (((-81.81169 24.56874 0.00000, ...
3  POLYGON Z ((-85.60516 34.98468 0.00000, -85.47...
4  POLYGON Z ((-117.24303 44.39097 0.00000, -117...
```

```
[28]: map_us = gp.read_file('usa-states-census-2014.shp')
      map_us.dtypes
```

```
[28]: STATEFP      object
      STATENS     object
      AFFGEOID    object
      GEOID       object
      STUSPS      object
      NAME        object
      LSAD        object
      ALAND       int64
      AWATER      int64
      region      object
      geometry    geometry
      dtype: object
```

```
[31]: map_us.join(df_copy)
```

```
bad_states = ['Commonwealth of the Northern Mariana Islands', 'Guam', 'United_
↳States Virgin Islands',
              'American Samoa', 'Puerto Rico', 'District of Columbia']
```

```

map_us_copy = map_us.copy()
map_us_copy = map_us[~map_us.NAME.isin(bad_states)]
map_us_copy = map_us_copy.sort_values(by='NAME')
new_map = pd.merge(map_us_copy, df_copy, left_on='STUSPS', right_on='State')

```

```
[37]: new_map = new_map.to_crs('EPSG:3395',);
```

```
[40]: new_map = new_map.round({'Partisan Lean': 2})
```

```
[41]: # Code adapted from https://jcutrer.com/python/learn-geopandas-plotting-usmaps
ax = new_map.boundary.plot(figsize=(18, 12), color='Black', linewidth=.5)
west = new_map[new_map['region'] == 'West']
southwest = new_map[new_map['region'] == 'Southwest']
southeast = new_map[new_map['region'] == 'Southeast']
midwest = new_map[new_map['region'] == 'Midwest']
northeast = new_map[new_map['region'] == 'Northeast']

west.apply(
    lambda x: ax.annotate(
        text="PL:" + str(x['Partisan Lean']),
        xy= (x.geometry.centroid.coords[0]),
        ha='center',
        color='#000000',
        fontsize=10),axis=1);

southwest.apply(
    lambda x: ax.annotate(
        text="PL:" + str(x['Partisan Lean']),
        xy= (x.geometry.centroid.coords[0]),
        ha='center',
        color='#000000',
        fontsize=10),axis=1);

southeast.apply(
    lambda x: ax.annotate(
        text="PL:" + str(x['Partisan Lean']),
        xy= (x.geometry.centroid.coords[0]),
        ha='center',
        color='#000000',
        fontsize=10),axis=1);

midwest.apply(
    lambda x: ax.annotate(
        text="PL:" + str(x['Partisan Lean']),
        xy= (x.geometry.centroid.coords[0]),
        ha='center',
        color='#000000',

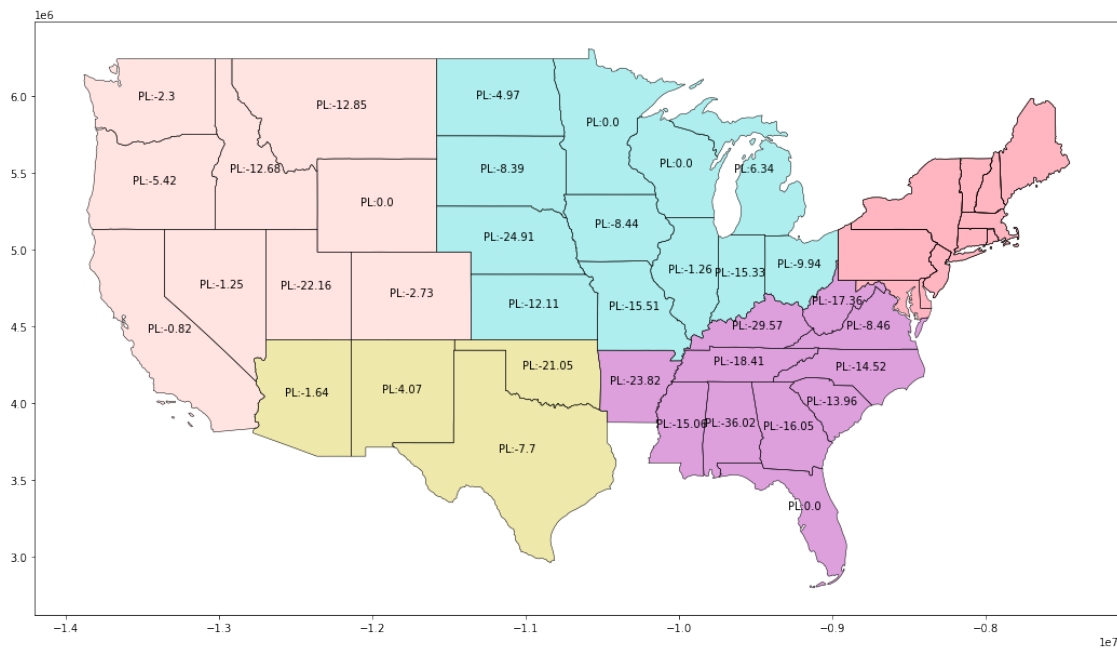
```

```

        fontsize=10),axis=1);

west.plot(ax=ax, color="MistyRose")
southwest.plot(ax=ax, color="PaleGoldenRod")
southeast.plot(ax=ax, color="Plum")
midwest.plot(ax=ax, color="PaleTurquoise")
final_map = northeast.plot(ax=ax, color="LightPink")

```



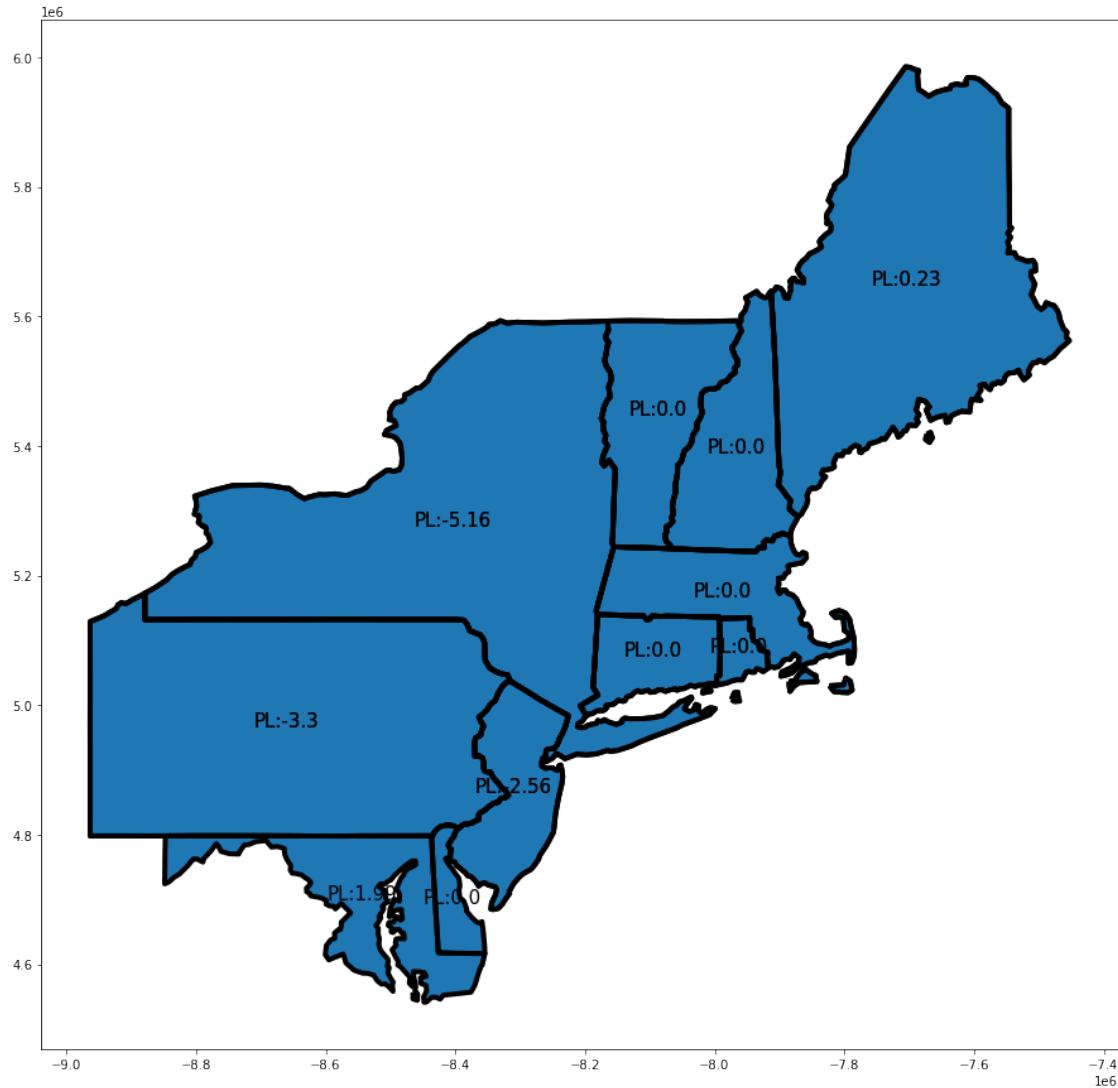
This map is a useful visualization of the partisan leans between states. We see that the South is heavily republican while the West is less republican and leaning towards more Democrat. This is quite as expected but can be useful when performing Bayesian Inference.

```

[43]: fig = plt.figure(1, figsize=(25,15))
      ax = fig.add_subplot()
      northeast = new_map[new_map['region'] == 'Northeast']
      northeast.apply(
          lambda x: ax.annotate(
              text="PL:" + str(x['Partisan Lean']),
              xy=(x.geometry.centroid.coords[0]),
              ha='center',
              color='#000000', # blue
              fontsize=15),axis=1);
      northeast.boundary.plot(ax=ax, color='Black', linewidth=4)
      northeast.plot(figsize=(10, 10), ax = ax)

```

[43]: <AxesSubplot:>



As expected with our prior beliefs the northeast seems to have a partisan lean more towards Democrats!

## 4 3.0 Nonparameteric Methods/GLMs.

Can we predict if a candidate won the primary based on the number of endorsements/type of endorsement

1. GLM : Logistic Regression because we have a binary decision.
2. Nonparametric Methods: DecisionTree/RandomForest because it again works well with binary decisions and our data should be easily seperable.

## 4.1 3.1 Nonparametric Methods.

### 4.1.1 3.1.1 Train Test Split

We will conduct a train test split with 70% of the data being in our training set.

```
[51]: # Train test split random seed 102
from sklearn.model_selection import train_test_split

train, test = train_test_split(df, train_size = 0.7, test_size = 0.3,
    random_state=10)
```

### 4.1.2 3.1.2 Feature Engineering

1. We will create a column of unique endorsements based on the number of endorsements from republicans and democrats.
2. Next we will calculate the proportion of endorsements based on party.

```
[52]: # Sum over all rows and find the number of unique endorsements.
train.loc[:, 'Unique Endorsements'] = train[endorsements].sum(axis = 1)
test.loc[:, 'Unique Endorsements'] = test[endorsements].sum(axis = 1)
```

/opt/conda/lib/python3.9/site-packages/pandas/core/indexing.py:1667:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

self.obj[key] = value

```
[53]: train.sample(n= 5)
```

```
[53]:
```

	Candidate	State		District	\
152	Paulette E. Jordan	ID		Governor of Idaho	
726	Pedro Valencia	TX		U.S. House Texas District 29	
459	Omar Vaid	NY		U.S. House New York District 11	
286	Janet T. Mills	ME		Governor of Maine	
362	David Wilson Brown	NC		U.S. House North Carolina District 10	

	Office Type	Race Type	Race	Primary Election Date	Primary Status	\
152	Governor	Regular		5/15/18	Advanced	
726	Representative	Regular		3/6/18	Lost	
459	Representative	Regular		6/26/18	Lost	
286	Governor	Regular		6/12/18	Advanced	
362	Representative	Regular		5/8/18	Advanced	

	Primary Runoff Status	General Status	Partisan Lean	...	NRA Endorsed?	\
152	None	On the Ballot	-34.330002	...	0	

726	None	None	39.970001 ...	0
459	None	None	-8.810000 ...	0
286	None	On the Ballot	3.510000 ...	0
362	None	On the Ballot	-25.510000 ...	0

	Right to Life Endorsed?	Susan B. Anthony Endorsed?	\
152	0	0	
726	0	0	
459	0	0	
286	0	0	
362	0	0	

	Club for Growth Endorsed?	Koch Support?	House Freedom Support?	\
152	0	0	0	
726	0	0	0	
459	0	0	0	
286	0	0	0	
362	0	0	0	

	Tea Party Endorsed?	Main Street Endorsed?	Chamber Endorsed?	\
152	0	0	0	
726	0	0	0	
459	0	0	0	
286	0	0	0	
362	0	0	0	

	Unique Endorsements
152	3
726	0
459	0
286	2
362	1

[5 rows x 47 columns]

```
[54]: def calculate_proportion(row):
      if (row['Party'] == 'Democrat'):
          row['Proportion'] = row['Unique Endorsements']/len(dem_endorsements)
      else:
          row['Proportion'] = row['Unique Endorsements']/len(rep_endorsements)
      return row
```

```
[55]: train = train.apply(lambda row: calculate_proportion(row), axis = 1)
      test = test.apply(lambda row: calculate_proportion(row), axis = 1)
```

```
[56]: X = train[['Partisan Lean', 'Unique Endorsements', 'Proportion', 'Party Support?
      ↪ ']]
```

```
y = train['Won Primary']
```

#### 4.1.3 3.1.3 Nonparametric Methods Results.

```
[182]: from sklearn.linear_model import LogisticRegression as LR

model = LR()
model.fit(X, y)
y_pred = model.predict(X)
y_pred_test = model.predict(test[['Partisan Lean', 'Unique Endorsements',
    ↳ 'Proportion', 'Party Support?']])

train['pred'] = y_pred
test['pred'] = y_pred_test

from sklearn.metrics import accuracy_score

train_accuracy = accuracy_score(train['pred'], train['Won Primary'])
test_accuracy = accuracy_score(test['pred'], test['Won Primary'])
train_accuracy, test_accuracy
```

```
[182]: (0.7376014427412083, 0.7016806722689075)
```

We don't have the best accuracy using logistic regression lets see if we can do better with random forest and perhaps our data is not very linearly separable.

```
[183]: # These are the columns we will use in our model.
X_col = ['Partisan Lean', 'Unique Endorsements', 'Proportion', 'Party Support?']
y_col = 'Won Primary'
```

```
[184]: # Predict and add these predictions to our dataframe
from sklearn.ensemble import RandomForestClassifier

tree_model = RandomForestClassifier()
X = train[X_col]
y = train[y_col]
tree_model.fit(X,y)
y_pred = tree_model.predict(X)
y_pred_test = tree_model.predict(test[X_col])

train['tree_pred'] = y_pred
test['tree_pred'] = y_pred_test

from sklearn.metrics import accuracy_score
```

```
train_accuracy = accuracy_score(train['tree_pred'], train['Won Primary'])
test_accuracy = accuracy_score(test['tree_pred'], test['Won Primary'])
train_accuracy, test_accuracy
```

[184]: (0.8151487826871056, 0.6932773109243697)

We will look at our mean squared error.

#### 4.1.4 3.1.4 Interpretation of results

For our first model logistic regression seemed to have trouble linearly separating the data thus resulting in around 70% accuracy. It

It seems that our training accuracy is not the most accurate at 81% training accuracy. Our test accuracy is even lower and this is most likely due to the fact that we are overfitting our data therefore compromising something in our model. To fix this, better feature selection/feature engineering can help overfitting and will be something to look at to fix for our final project.

### 4.2 3.2 GLM Model

Next we will look at using a GLM logistic regression model. We will approach this using a frequentist approach then look at it from a bayesian approach.

```
[225]: from scipy.stats import poisson, norm, gamma
import statsmodels.api as sm

try:
    from pymc3 import *
    import pymc3 as pm
except:
    ! pip install pymc3
    from pymc3 import *
    import pymc3 as pm
```

```
[57]: # Some glm preprocessing
glm_df = train.copy()
glm_df['Unique Endorsements'] = glm_df[endorsements].sum(axis = 1)
```

#### 4.3 3.2.1 Methods

Again we are trying to predict the same thing as above in our nonparametric methods. We will be using a binomial glm because we have count data and our link function should work for our binary data. We will then evaluate the model by first using a frequentist model and analyzing the deviance and Pearson chi-squared before than analyzing the posterior distribution.



#### 4.4 3.2.2 Results

```
[18]: # Frequentist GLM
glm_binom = sm.GLM(glm_df['Won Primary'], exog=sm.add_constant(glm_df[X_col]),
                    family=sm.families.Binomial())
glm_res = glm_binom.fit()
print(glm_res.summary())
```

```
Generalized Linear Model Regression Results
=====
Dep. Variable:          Won Primary    No. Observations:          1109
Model:                  GLM           Df Residuals:              1105
Model Family:           Binomial      Df Model:                  3
Link Function:          logit         Scale:                    1.0000
Method:                 IRLS          Log-Likelihood:            -637.71
Date:                   Thu, 09 Dec 2021 Deviance:                  1275.4
Time:                   15:02:35       Pearson chi2:              1.11e+03
No. Iterations:         4
Covariance Type:        nonrobust
=====
=====
                    coef      std err          z      P>|z|      [0.025
0.975]
-----
const             -1.3622      0.096    -14.119      0.000     -1.551
-1.173
Partisan Lean     -0.0150      0.004     -3.542      0.000     -0.023
-0.007
Unique Endorsements -0.5497      0.265     -2.074      0.038     -1.069
-0.030
Proportion        16.4536      4.596      3.580      0.000      7.445
25.462
=====
=====
```

```
/opt/conda/lib/python3.9/site-packages/statsmodels/tsa/tsatools.py:142:
FutureWarning: In a future version of pandas all arguments of concat except for
the argument 'objs' will be keyword-only
    x = pd.concat(x[:,order], 1)
```

```
[19]: glm_df = glm_df.rename(columns={'Won Primary' : 'Won_Primary', 'Unique_
↳Endorsements' : 'Unique_Endorsements',
                                   'Partisan Lean': 'Partisan_Lean'})
```

```
[20]: # Bayesian GLM
with pm.Model() as bayes_model:
```

```

glm.GLM.from_formula("Won_Primary ~ + Unique_Endorsements + Partisan_Lean +
↪Proportion", glm_df,
                        family = glm.families.Binomial())
trace_binomial = pm.sample(1000, cores = 1, target_accept = 0.95,
↪init='adapt_diag')

```

/tmp/ipykernel\_24/536276188.py:5: FutureWarning: In v4.0, pm.sample will return an `arviz.InferenceData` object instead of a `MultiTrace` by default. You can pass return\_inferencedata=True or return\_inferencedata=False to be safe and silence this warning.

```

trace_binomial = pm.sample(1000, cores = 1, target_accept = 0.95,
init='adapt_diag')

```

Auto-assigning NUTS sampler...

Initializing NUTS using adapt\_diag...

Sequential sampling (2 chains in 1 job)

NUTS: [Proportion, Partisan\_Lean, Unique\_Endorsements, Intercept]

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Sampling 2 chains for 1\_000 tune and 1\_000 draw iterations (2\_000 + 2\_000 draws total) took 89 seconds.

```
[21]: pm.summary(trace_binomial)
```

Got error No model on context stack. trying to find log\_likelihood in translation.

/opt/conda/lib/python3.9/site-packages/arviz/data/io\_pymc3\_3x.py:98:

FutureWarning: Using `from\_pymc3` without the model will be deprecated in a future release. Not using the model will return less accurate and less useful results. Make sure you use the model argument or call from\_pymc3 within a model context.

```
warnings.warn(
```

```
[21]:
```

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	\
Intercept	-1.371	0.095	-1.562	-1.207	0.003	0.002	
Unique_Endorsements	-0.576	0.256	-1.002	-0.073	0.009	0.006	
Partisan_Lean	-0.015	0.004	-0.023	-0.007	0.000	0.000	
Proportion	16.953	4.430	8.326	24.704	0.158	0.113	

	ess_bulk	ess_tail	r_hat
Intercept	1027.0	1008.0	1.0
Unique_Endorsements	800.0	831.0	1.0
Partisan_Lean	969.0	983.0	1.0
Proportion	780.0	775.0	1.0

```
[22]: az.plot_trace(trace_binomial)
```

```

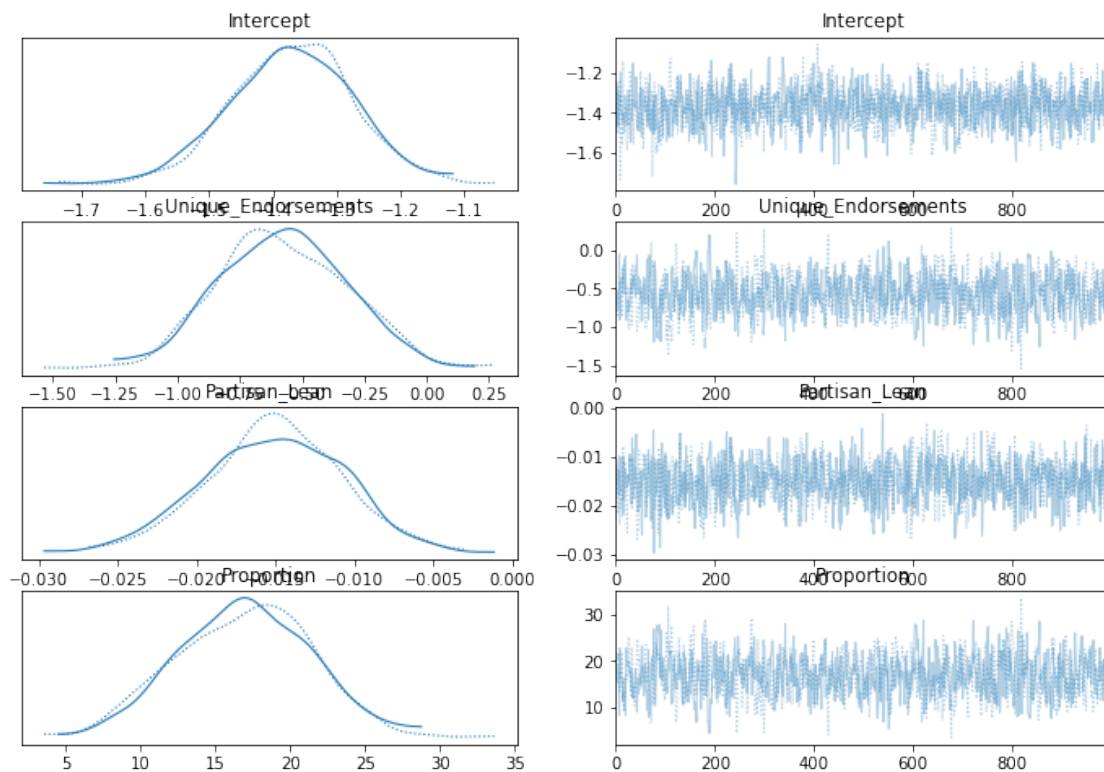
Got error No model on context stack. trying to find log_likelihood in
translation.
/opt/conda/lib/python3.9/site-packages/arviz/data/io_pymc3_3x.py:98:
FutureWarning: Using `from_pymc3` without the model will be deprecated in a
future release. Not using the model will return less accurate and less useful
results. Make sure you use the model argument or call from_pymc3 within a model
context.
    warnings.warn(
Got error No model on context stack. trying to find log_likelihood in
translation.

```

```

[22]: array([[<AxesSubplot:title={'center':'Intercept'}>,
               <AxesSubplot:title={'center':'Intercept'}>],
             [<AxesSubplot:title={'center':'Unique_Endorsements'}>,
               <AxesSubplot:title={'center':'Unique_Endorsements'}>],
             [<AxesSubplot:title={'center':'Partisan_Lean'}>,
               <AxesSubplot:title={'center':'Partisan_Lean'}>],
             [<AxesSubplot:title={'center':'Proportion'}>,
               <AxesSubplot:title={'center':'Proportion'}>]], dtype=object)

```



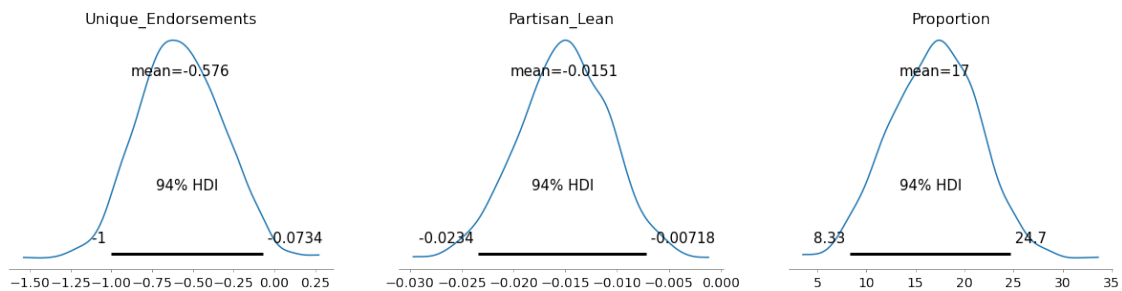
We can also plot the posterior distributions and see their credible intervals.

```
[23]: # Plot Posterior
az.plot_posterior(trace_binomial, ['Unique_Endorsements', 'Partisan_Lean', 'Proportion'], round_to=3)
```

Got error No model on context stack. trying to find log\_likelihood in translation.  
/opt/conda/lib/python3.9/site-packages/arviz/data/io\_pymc3\_3x.py:98:  
FutureWarning: Using `from\_pymc3` without the model will be deprecated in a future release. Not using the model will return less accurate and less useful results. Make sure you use the model argument or call from\_pymc3 within a model context.

```
warnings.warn(
```

```
[23]: array([<AxesSubplot:title={'center':'Unique_Endorsements'}>,
<AxesSubplot:title={'center':'Partisan_Lean'}>,
<AxesSubplot:title={'center':'Proportion'}>], dtype=object)
```



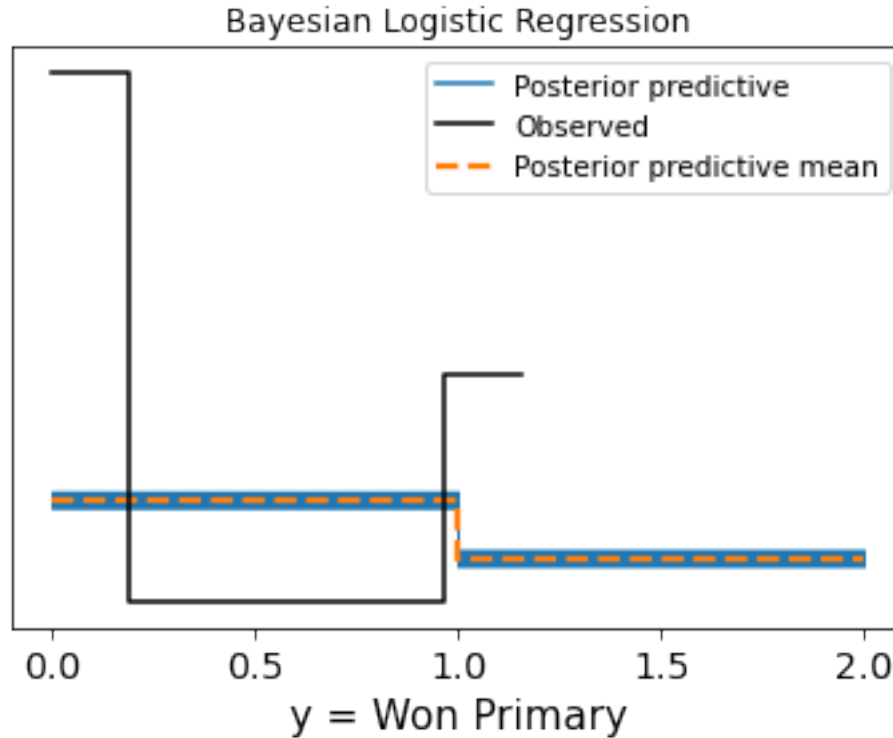
```
[61]: glm_df['Unique_Endorsements'].mean(), glm_df['Partisan_Lean'].mean(),
glm_df['Proportion'].mean()
```

```
[61]: (0.8917944093778178, -7.363444505509468, 0.05204017187311762)
```

```
[24]: # Bayesian Posterior Predictive Check
# THERE ARE SOME ISSUES HERE IDK WHY
with bayes_model:
    binomial_ppc = pm.sample_posterior_predictive(trace_binomial)
    ppc_binomial = az.from_pymc3(trace_binomial,
posterior_predictive=binomial_ppc)

az.plot_ppc(ppc_binomial)
plt.xlabel('y = Won Primary')
plt.title('Bayesian Logistic Regression')
plt.show()
```

```
<IPython.core.display.HTML object>
```



#### 4.4.1 3.2.3 Results for GLM and Comparison to Random Forest

For our frequentist model we see we have a very negative log likelihood which may seem that this is not a good fit for our model. However if we look at the number of observations minus the number of parameters  $n - p = 1109 - 5 = 1104$ . Our deviance and chi-square  $\chi^2$  is 1275 and 1100 respectively because these are not too off our frequentist model does not seem to be a very good fit.

Comparing to our nonparametric method of random forest we found we had a not terrible training accuracy but not very good test accuracy. We concluded this must be because of overfitting. However our GLM shows that indeed our model has some problems with our frequentist as well as our bayesian model using posterior predictive checks.

```
[38]: poll_data = pd.read_csv('presidential_polls.csv')
```

```
[64]: poll_data.head(2)
```

```
[64]:
```

	startdate	enddate	state	pollster	grade	poll_wt	\
0	2016-10-25	2016-10-31	U.S.	Google Consumer Surveys	B	6.139129	
1	2016-10-27	2016-10-30	U.S.	ABC News/Washington Post	A+	4.197292	

	samplesize	population	adjpoll_clinton	adjpoll_trump
0	24316.0	lv	42.64140	40.86509
1	1128.0	lv	43.29659	44.72984

```
[40]: poll_data.columns
```

```
[40]: Index(['cycle', 'branch', 'type', 'matchup', 'forecastdate', 'state',
        'startdate', 'enddate', 'pollster', 'grade', 'samplesize', 'population',
        'poll_wt', 'rawpoll_clinton', 'rawpoll_trump', 'rawpoll_johnson',
        'rawpoll_mcmullin', 'adjpoll_clinton', 'adjpoll_trump',
        'adjpoll_johnson', 'adjpoll_mcmullin', 'multiversions', 'url',
        'poll_id', 'question_id', 'createddate', 'timestamp'],
        dtype='object')
```

```
[ ]:
```

## 4.5 4.0 EDA on Polling Data from Kaggle.

We will begin with some EDA to visualize if we can identify key features or plot data to help form good priors.

### 4.5.1 Dataset was taken from <https://www.kaggle.com/fivethirtyeight/2016-election-polls>

```
[68]: poll_data = pd.read_csv('presidential_polls.csv')
```

```
[69]: poll_data.head(2)
```

```
[69]:
```

	cycle	branch	type	matchup	forecastdate	\
0	2016	President	polls-plus	Clinton vs. Trump vs. Johnson	11/1/16	
1	2016	President	polls-plus	Clinton vs. Trump vs. Johnson	11/1/16	

	state	startdate	enddate	pollster	grade	...	\
0	U.S.	10/25/2016	10/31/2016	Google Consumer Surveys	B	...	
1	U.S.	10/27/2016	10/30/2016	ABC News/Washington Post	A+	...	

	adjpoll_clinton	adjpoll_trump	adjpoll_johnson	adjpoll_mcmullin	\
0	42.64140	40.86509	5.675099	NaN	
1	43.29659	44.72984	3.401513	NaN	

	multiversions	url	poll_id	\
0	NaN	<a href="https://datastudio.google.com/u/0/#/org//repor...">https://datastudio.google.com/u/0/#/org//repor...</a>	47940	
1	NaN	<a href="http://www.langerresearch.com/wp-content/uploa...">http://www.langerresearch.com/wp-content/uploa...</a>	47881	

	question_id	createddate	timestamp	
0	74999	11/1/16	15:09:38	1 Nov 2016
1	74936	11/1/16	15:09:38	1 Nov 2016

```
[2 rows x 27 columns]
```

```
[70]: poll_data.columns
```

```
[70]: Index(['cycle', 'branch', 'type', 'matchup', 'forecastdate', 'state',
          'startdate', 'enddate', 'pollster', 'grade', 'samplesize', 'population',
          'poll_wt', 'rawpoll_clinton', 'rawpoll_trump', 'rawpoll_johnson',
          'rawpoll_mcmullin', 'adjpoll_clinton', 'adjpoll_trump',
          'adjpoll_johnson', 'adjpoll_mcmullin', 'multiversions', 'url',
          'poll_id', 'question_id', 'createddate', 'timestamp'],
          dtype='object')
```

```
[71]: poll_data = poll_data[["startdate", "enddate", "state", "pollster", "grade",
    ↪ 'poll_wt',
    "samplesize", "population", "adjpoll_clinton", "adjpoll_trump"]]

poll_data['startdate'] = pd.to_datetime(poll_data["startdate"])
poll_data["enddate"] = pd.to_datetime(poll_data["enddate"])
poll_data.sort_values(by=['startdate', 'enddate'])
```

```
[71]:
```

	startdate	enddate	state \
3071	2015-11-06	2015-11-16	New Hampshire
6517	2015-11-06	2015-11-16	New Hampshire
9895	2015-11-06	2015-11-16	New Hampshire
3077	2015-11-07	2015-11-08	South Carolina
6520	2015-11-07	2015-11-08	South Carolina
...	...	...	...
6864	2016-10-30	2016-10-30	Virginia
6872	2016-10-30	2016-10-30	Florida
6881	2016-10-30	2016-10-30	Colorado
6905	2016-10-30	2016-10-30	Nevada
6987	2016-10-30	2016-10-30	Michigan

	pollster	grade	poll_wt	samplesize \
3071	Morning Consult	NaN	6.120000e-06	530.0
6517	Morning Consult	NaN	7.660000e-07	530.0
9895	Morning Consult	NaN	6.120000e-06	530.0
3077	Public Policy Polling	B+	5.460000e-06	1290.0
6520	Public Policy Polling	B+	6.710000e-07	1290.0
...	...	...	...	...
6864	Remington	NaN	2.282895e+00	1106.0
6872	Remington	NaN	2.171098e+00	989.0
6881	Remington	NaN	2.132726e+00	952.0
6905	Remington	NaN	1.940524e+00	787.0
6987	Mitchell Research & Communications	D	1.417711e+00	953.0

	population	adjpoll_clinton	adjpoll_trump
3071	rv	45.40125	38.46687
6517	rv	45.20979	38.59373
9895	rv	45.39273	38.40193

3077	v	40.63694	45.19826
6520	v	40.48820	45.19637
...	...	...	...
6864	lv	46.66191	41.12267
6872	lv	43.66191	46.12267
6881	lv	44.66191	42.12267
6905	lv	43.66191	46.12267
6987	lv	45.61794	40.12326

[10236 rows x 10 columns]

Next we will identify battleground states based on the difference between the adjusted poll values between Trump and Clinton. If the difference between clinton and the trump poll fall below a value  $\epsilon$  we will define this as a battleground state.

```
[72]: df = poll_data.copy()
```

```
[73]: 57.9 + 33
```

```
[73]: 90.9
```

```
[74]: df_group = df.groupby(by=['state']).mean().reset_index()
df_group.head(5)
```

```
[74]:
```

	state	poll_wt	samplesize	adjpoll_clinton	adjpoll_trump
0	Alabama	0.106648	608.968750	32.570923	57.944145
1	Alaska	0.263570	319.310345	35.678765	43.713035
2	Arizona	0.305452	760.866667	41.697958	43.768535
3	Arkansas	0.143535	335.882353	37.011808	50.654389
4	California	0.224925	1275.036364	54.934097	32.294801

```
[75]: def battleground(row, eps=2):
    if abs(row['adjpoll_clinton'] - row['adjpoll_trump']) < eps:
        row['battleground'] = 1
    else:
        row['battleground'] = 0
    return row
```

```
[76]: df_group = df_group.apply(lambda row: battleground(row), axis = 1)
```

```
[77]: battleground_states = df_group[df_group['battleground'] == 1]['state']

df[df['state'].isin(battleground_states.tolist())]
battleground_states
```

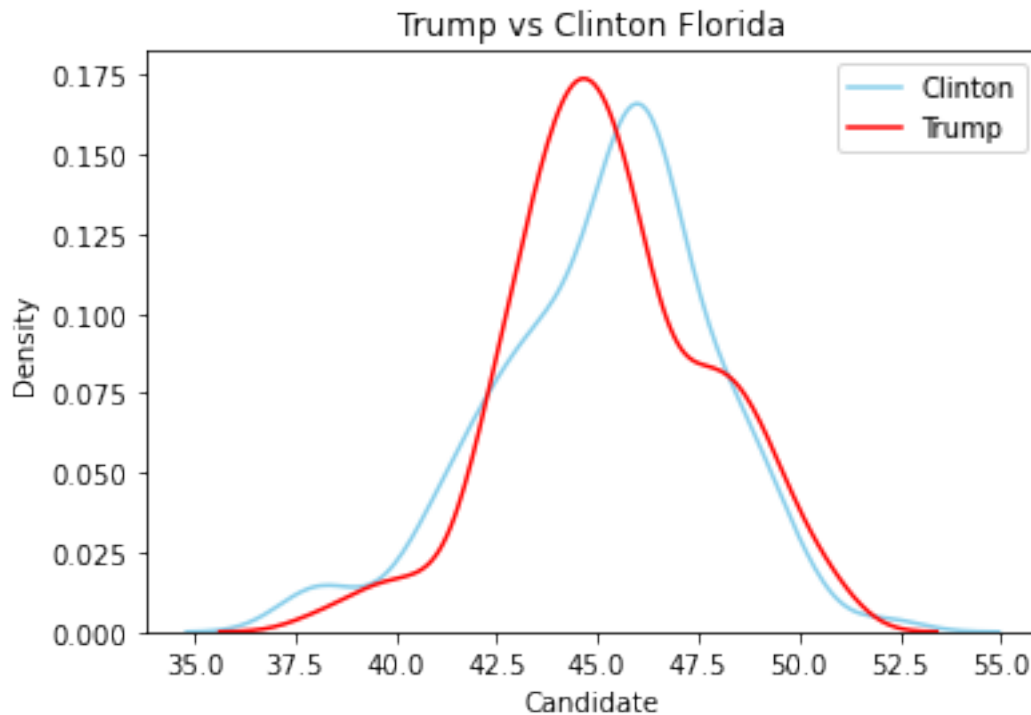
```
[77]: 9          Florida
      15          Iowa
```



```
33         Nevada
38     North Carolina
40         Ohio
Name: state, dtype: object
```

```
[78]: # Look at florida
fig, ax = plt.subplots()
sns.kdeplot(df[df['state'] == 'Florida']['adjpoll_clinton'], ax = ax,
            color='skyblue', label='Clinton')
sns.kdeplot(df[df['state'] == 'Florida']['adjpoll_trump'], ax = ax, color =
            'red', label='Trump')
plt.xlabel("Candidate")
plt.title("Trump vs Clinton Florida")
plt.legend()
```

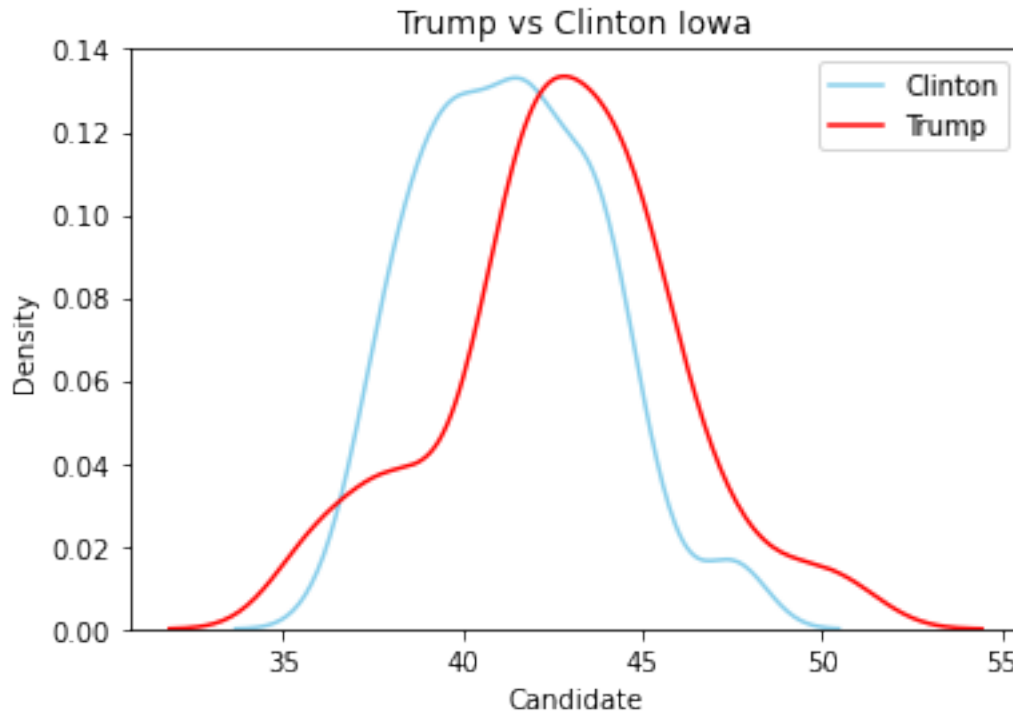
```
[78]: <matplotlib.legend.Legend at 0x7efb7c0a0c40>
```



```
[79]: fig, ax = plt.subplots()
sns.kdeplot(df[df['state'] == 'Iowa']['adjpoll_clinton'], ax = ax,
            color='skyblue', label='Clinton')
sns.kdeplot(df[df['state'] == 'Iowa']['adjpoll_trump'], ax = ax, color = 'red',
            label='Trump')
plt.xlabel('Candidate')
```

```
plt.title('Trump vs Clinton Iowa')
plt.legend()
```

[79]: <matplotlib.legend.Legend at 0x7efb6cc8edf0>



We can see for battleground states even though a state has a higher mode such as Clinton in Florida, Trump in the polls still has the lead as his average poll rating is higher. This is most likely due to that Trump has a higher density in Florida giving him a slight lead. Most of the differences come from the distributions being slightly wider or taller in key areas.

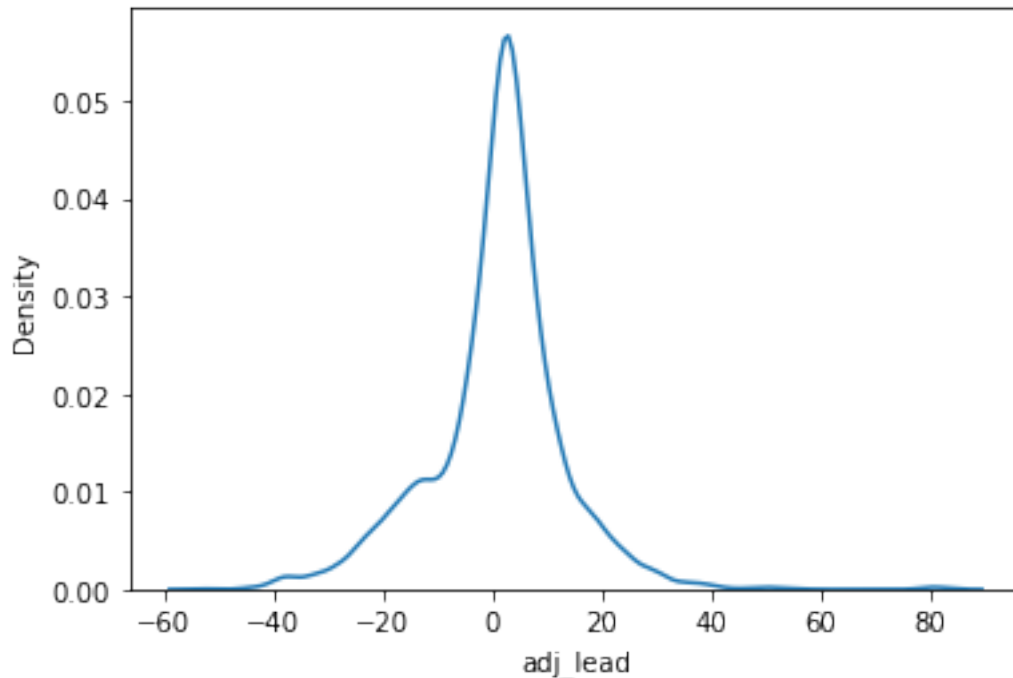
Lets introduce a new variable adjusted lead which will be the difference between Clinton and Trump poll ratings where a positive value means Clinton has a lead over trump and negative means Trump has the lead over Clinton.

```
[80]: df['adj_lead'] = df['adjpoll_clinton'] - df['adjpoll_trump']
      df['adj_lead'].mean(), df['adj_lead'].var()
```

[80]: (0.9762825313599062, 153.4731635576654)

```
[81]: sns.kdeplot(df['adj_lead'])
```

[81]: <AxesSubplot:xlabel='adj\_lead', ylabel='Density'>



We are trying to construct a prior belief on our estimate for the adjusted lead of the candidate. We will assume that our likelihood follows a normal distribution.

From here we can see the partisan leans for each state and which ones are leaning red or blue. We can see heavy republican leans on the southeast while on the west it seems that leans are much

## 5 5.0 Bayesian Hierarchial Modeling

**5.0.1 Our resarch question will be dealing with trying to estimate the poll leads between Clinton and Trump during the 2016 presidential election. Dataset was taken from <https://www.kaggle.com/fivethirtyeight/2016-election-polls>**

### 5.1 5.1 Building a Basic Model/Methods

Our graphical model is as follows whre we have our parameter where we have this as our prior belief based on our estimates above.

$$\mu \sim Uniform(-40, 40)$$

$$\sigma^2 \sim HalfNormal(0, 5)$$

$$X \mid (\mu, \sigma^2) \sim \mathcal{N}(\mu, \sigma^2)$$

We choose  $\mu$  and  $\sigma^2$  as these priors because we see that most of the values fall between -60 and 60. Finally the variance cannot be negative and the we have a tight distribution so we choose a smaller variance so we choose a half normal.

```
[82]: import pymc3 as pm
```

```
[83]: with pm.Model() as model:
        = pm.Uniform(' ', lower=-60, upper=60)
        = pm.HalfNormal(' ', sd=5)
        X = pm.Normal('X', mu= , sd= , observed=df['adj_lead'].values)
        trace_guassian = pm.sample(1000, tune=1000, return_inferencedata=True)
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt\_diag...

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [ , ]

<IPython.core.display.HTML object>

Sampling 4 chains for 1\_000 tune and 1\_000 draw iterations (4\_000 + 4\_000 draws total) took 6 seconds.

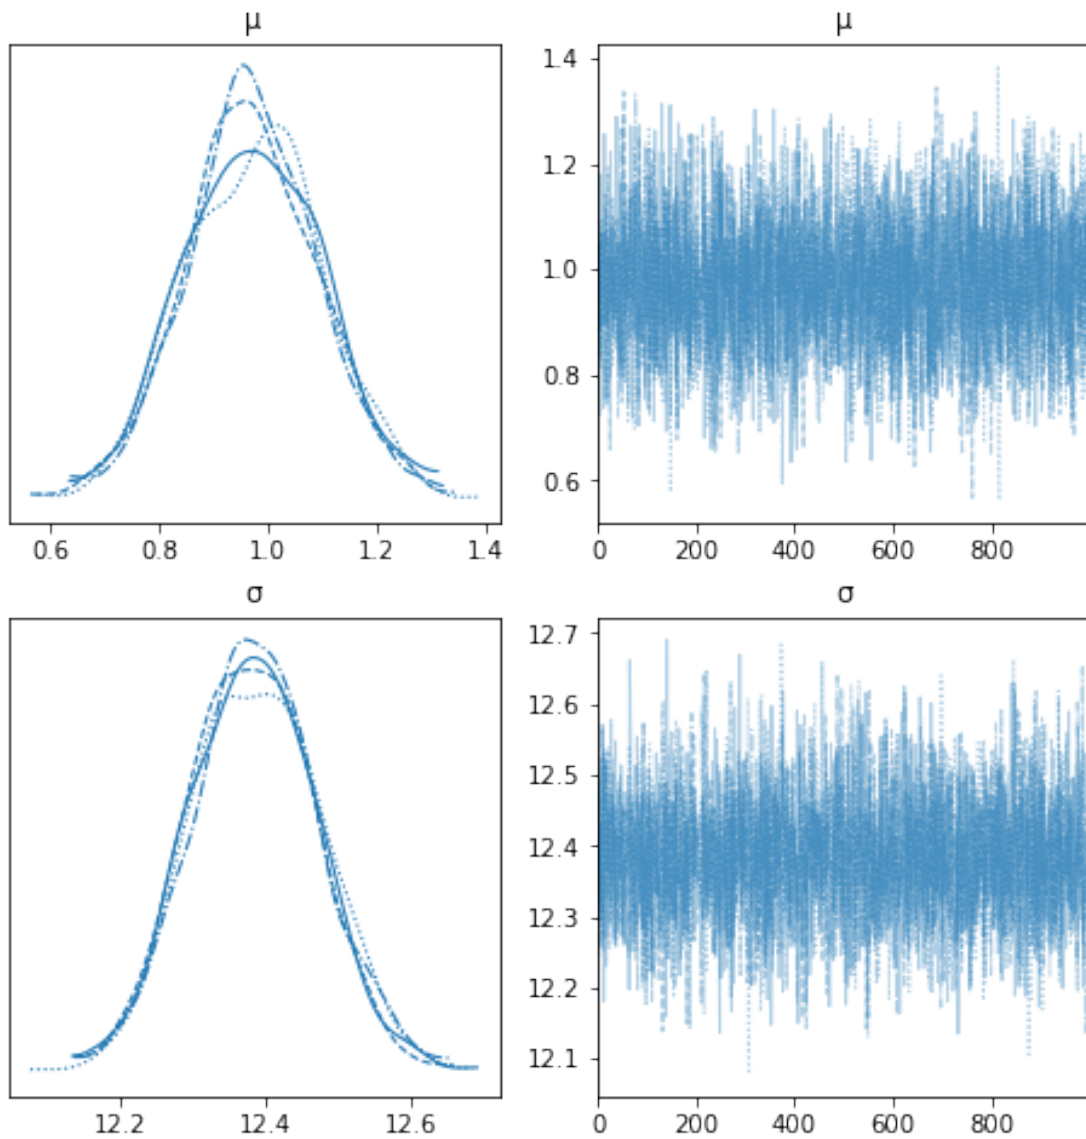
## 5.2 5.2 Results

```
[85]: import arviz as az

f, axs = plt.subplots(2,2,figsize=(8,8))

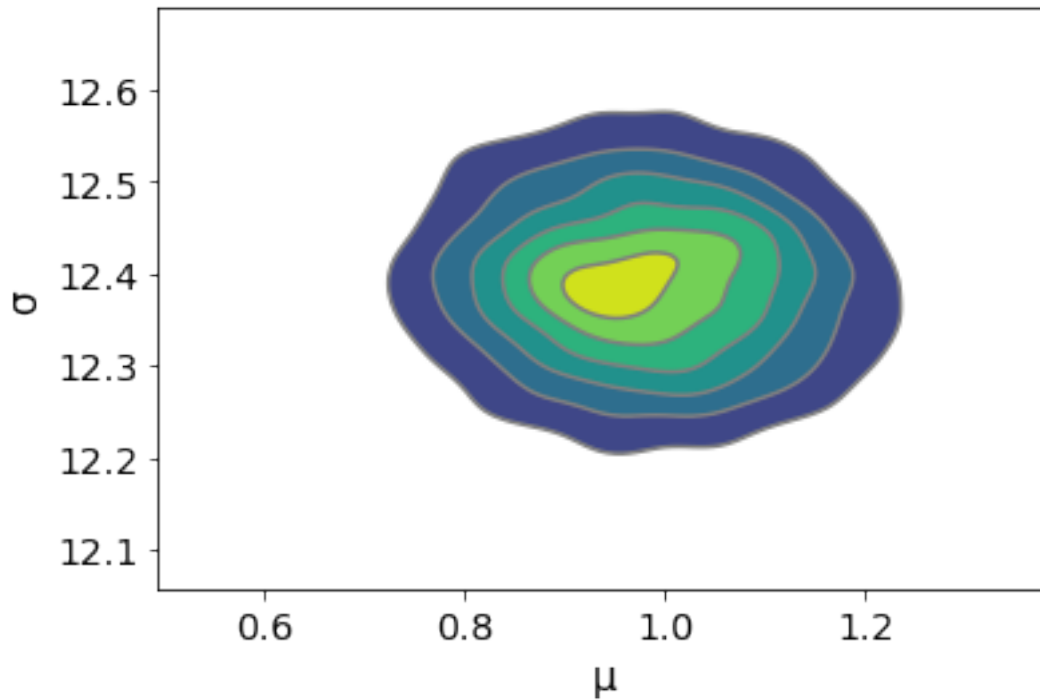
az.plot_trace(trace_guassian, axes = axs)
```

```
[85]: array([[<AxesSubplot:title={'center':' '}>,
             <AxesSubplot:title={'center':' '}>],
            [<AxesSubplot:title={'center':' '}>,
             <AxesSubplot:title={'center':' '}>]], dtype=object)
```



Our trace plots look good here. Our values on the left seem to have converged and be stationary and our MAP estimate which is the peak on the left graphs seem to be relatively close to the true value. Lets look at the joint trace of our sample.

```
[21]: az.plot_pair(trace_gaussian, kind='kde', fill_last=False);
```



We can see here that our values do not seem to be correlated with each other which is a good thing. Now lets see if we can categorize the states into groups.

### 5.3 Adding Regions to our Model

Lets consider adding the four regions to our model and consider how this changes our model.

```
[20]: ne = ['Connecticut', 'Maine', 'Massachusetts', 'New Hampshire', 'Rhode Island',
           'Vermont', 'New Jersey', 'New York', 'Pennsylvania', 'Maine CD-1', 'Maine
           ↳CD-2']

mw = ['Illinois', 'Indiana', 'Michigan', 'Ohio', 'Wisconsin', 'Iowa', 'Kansas', '
           ↳Minnesota', 'Missouri',
           'Nebraska', 'North Dakota', 'South Dakota', 'Nebraska CD-2', 'Nebraska
           ↳CD-3', 'Nebraska CD-1']

so= ['Delaware', 'Florida', 'Georgia', 'Maryland', 'North Carolina',
     'South Carolina', 'Virginia', 'District of Columbia', 'West Virginia',
     'Alabama', 'Kentucky', 'Mississippi', 'Tennessee', 'Arkansas', 'Louisiana',
     ↳'Oklahoma', 'Texas']

we = ['Arizona', 'Colorado', 'Idaho', 'Montana', 'Nevada', 'New Mexico',
     ↳'Utah', 'Wyoming',
     'Alaska', 'California', 'Hawaii', 'Oregon', 'Washington']
```

```
[97]: # U.S randomly considred a state.
df_2 = df.copy()
df_2 = df[df['state'] != 'U.S.']
```

```
[99]: # Adds a region column to each row which is encoded.
def add_region(row):
    if row['state'] in ne:
        row['region'] = 0
    elif row['state'] in mw:
        row['region'] = 1
    elif row['state'] in so:
        row['region'] = 2
    elif row['state'] in we:
        row['region'] = 3
    return row
```

```
[98]: # Apply our function we created above.
df_2 = df_2.apply(lambda row: add_region(row), axis= 1)
```

```
[100]: # Create a index to use in our model.
idx = pd.Categorical(df_2['region'],
                    categories=[0, 1, 2, 3]).codes
```

## 5.4 5.4 Results of More Robust Model

```
[32]: with pm.Model() as model_regions:
        = pm.Uniform(' ', lower=-60, upper=60,shape= 4)
        = pm.HalfNormal(' ', sd=5,shape=4)
        X = pm.Normal('X', mu= [idx], sd= [idx], observed=df_2['adj_lead'].values)
        trace_guassian_regions = pm.sample(1000, tune=1000,
        ↪return_inferencedata=False, cores = 1)
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt\_diag...

Sequential sampling (2 chains in 1 job)

NUTS: [ , ]

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Sampling 2 chains for 1\_000 tune and 1\_000 draw iterations (2\_000 + 2\_000 draws total) took 14 seconds.

```
[34]: az.summary(az.from_pymc3(trace_guassian_regions))
```

Got error No model on context stack. trying to find log\_likelihood in translation.

```
[34]:
```

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	\
[0]	10.408	0.268	9.906	10.878	0.005	0.003	3228.0	1453.0	
[1]	-2.038	0.268	-2.507	-1.498	0.005	0.004	2580.0	1183.0	
[2]	-5.155	0.318	-5.752	-4.557	0.006	0.004	2828.0	1611.0	
[3]	0.661	0.377	-0.083	1.352	0.006	0.005	3789.0	1608.0	
[0]	9.759	0.186	9.406	10.084	0.004	0.002	2790.0	1768.0	
[1]	11.088	0.188	10.749	11.456	0.003	0.002	4227.0	1657.0	
[2]	15.646	0.227	15.194	16.050	0.004	0.003	3370.0	1694.0	
[3]	14.614	0.249	14.189	15.110	0.004	0.003	3630.0	1770.0	

```

r_hat
[0] 1.0
[1] 1.0
[2] 1.0
[3] 1.0
[0] 1.0
[1] 1.0
[2] 1.0
[3] 1.0

```

From this summary table we can see that there is quite a different in regions. For example the north east seems to have a big lead towards Hilary Clinton while the midwest and south favor Trump much more.

```
[77]: az.plot_trace(trace_guassian_regions)
```

```

Got error No model on context stack. trying to find log_likelihood in
translation.
/opt/conda/lib/python3.9/site-packages/arviz/data/io_pymc3_3x.py:98:
FutureWarning: Using `from_pymc3` without the model will be deprecated in a
future release. Not using the model will return less accurate and less useful
results. Make sure you use the model argument or call from_pymc3 within a model
context.

```

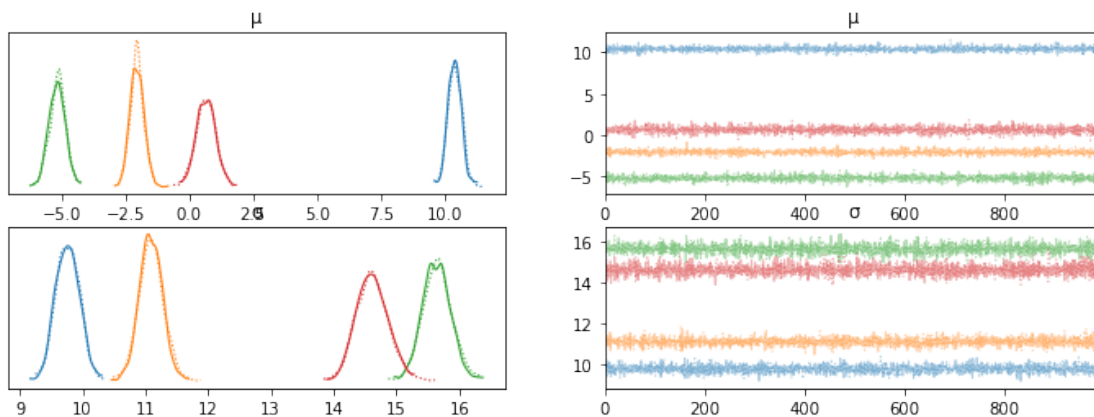
```

warnings.warn(
Got error No model on context stack. trying to find log_likelihood in
translation.

```

```
[77]: array([[<AxesSubplot:title={'center':' '>,
<AxesSubplot:title={'center':' '>],
[<AxesSubplot:title={'center':' '>,
<AxesSubplot:title={'center':' '>]], dtype=object)
```





```
[78]: df_2.groupby('region').mean()
```

```
[78]:
```

	samplesize	adjpoll_clinton	adjpoll_trump	adj_lead
region				
0	634.572368	47.855455	37.453344	10.402111
1	714.257391	40.852794	42.884180	-2.031386
2	675.943750	41.185942	46.355735	-5.169793
3	606.754128	41.262549	40.590502	0.672047

We see that our model does a good job of being accurate in estimating the means between region. Lets next run a posterior predictive check as model checking is always good.

```
[82]: ppc = pm.sample_posterior_predictive(trace_guassian_regions, samples=2000,
↳ model=model_regions)
```

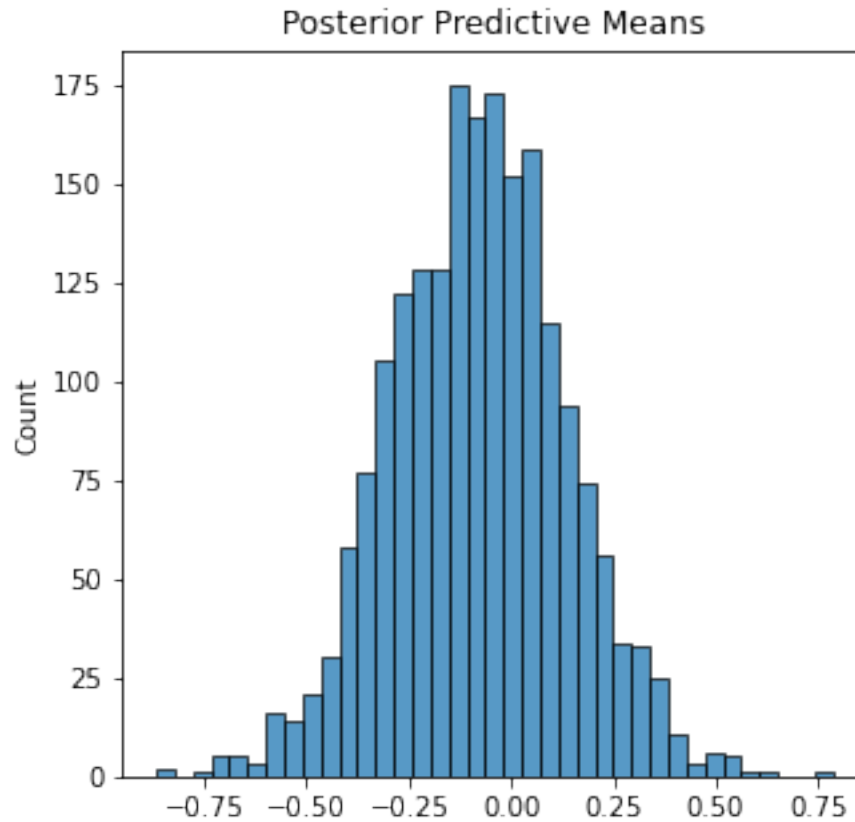
<IPython.core.display.HTML object>

```
[84]: np.asarray(ppc['X']).shape
```

```
[84]: (2000, 7131)
```

```
[91]: means = [X.mean() for X in ppc['X']]
ax = plt.subplots(figsize=(5, 5))
sns.histplot(means)
plt.title('Posterior Predictive Means')
```

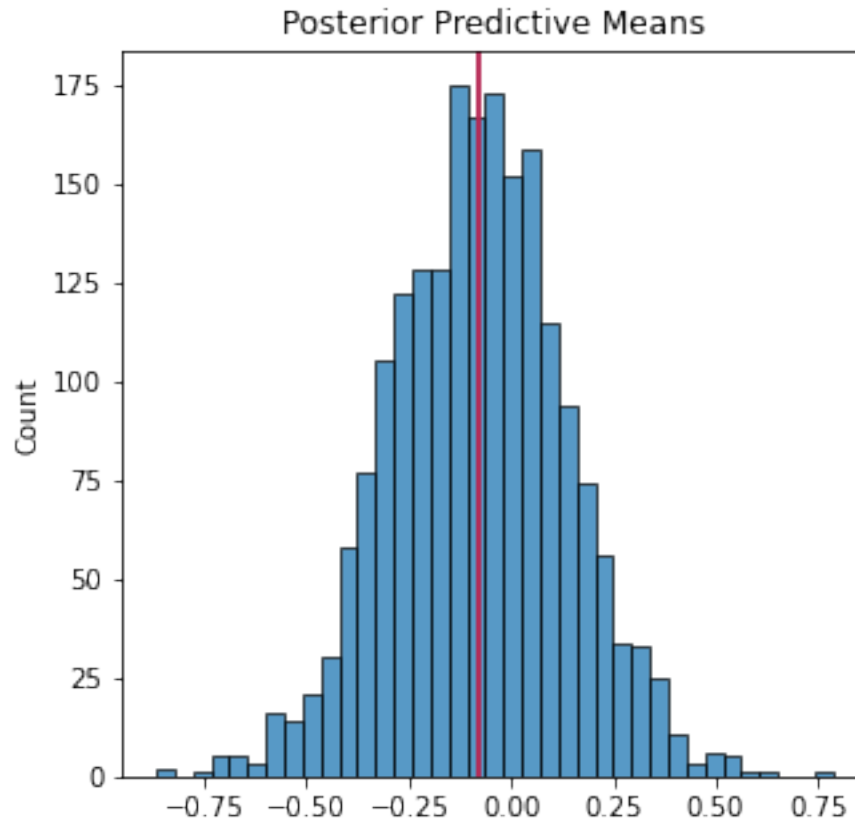
```
[91]: Text(0.5, 1.0, 'Posterior Predictive Means')
```



Lets overlay a line of the true mean.

```
[96]: means = [X.mean() for X in ppc['X']]
ax = plt.subplots(figsize=(5, 5))
sns.histplot(means)
plt.title('Posterior Predictive Means')
ax[1].axvline(df_2['adj_lead'].mean())
plt.axvline(df_2['adj_lead'].mean(), color='crimson')
```

```
[96]: <matplotlib.lines.Line2D at 0x7fd03efcbb50>
```



We see here that our model does a good job of approximating the mean!

## 5.5 5.4 Discussion

We began with a simple model where we used our prior beliefs and gaussian likelihood function to approximate the adjusted lead which was the difference between the Trump and Clinton poll. We then tried a more robust graphical model where we introduced regions and from there saw differences in adjusted leads based on the region. I tried to incorporate poll weight into our graphical model but ran into trouble as I didn't know how the likelihood would change or maybe the prior I chose for the poll weight was incorrect. This is one of the issues with Bayesian Inference is that you don't quite know how to choose a prior which was the issue I ran into. Some data that would be interesting to look at were the demographics in each region such as gender, race, etc.

[ ]: